

# *Siddhi-CEP*

## **A SECOND LOOK AT COMPLEX EVENT PROCESSING**

Srinath Perera,  
Senior Software Architect WSO2 Inc.  
Visiting Faculty, University of Moratuwa  
Member, Apache Software Foundation

<http://siddhi.sourceforge.net/>



## OUTLINE

- Introduction to CEP
- Second look at CEP Implementations
- Siddhi Architecture
- Siddhi Performance
- Conclusion and Future Topics

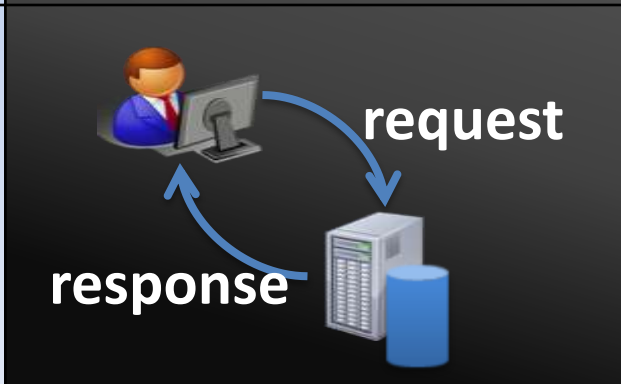
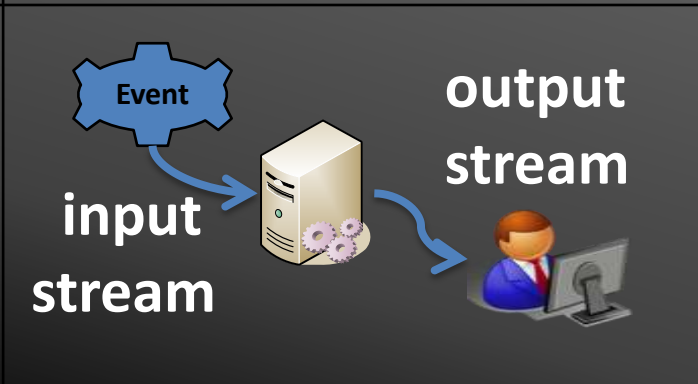


photo by John Trainoron Flickr

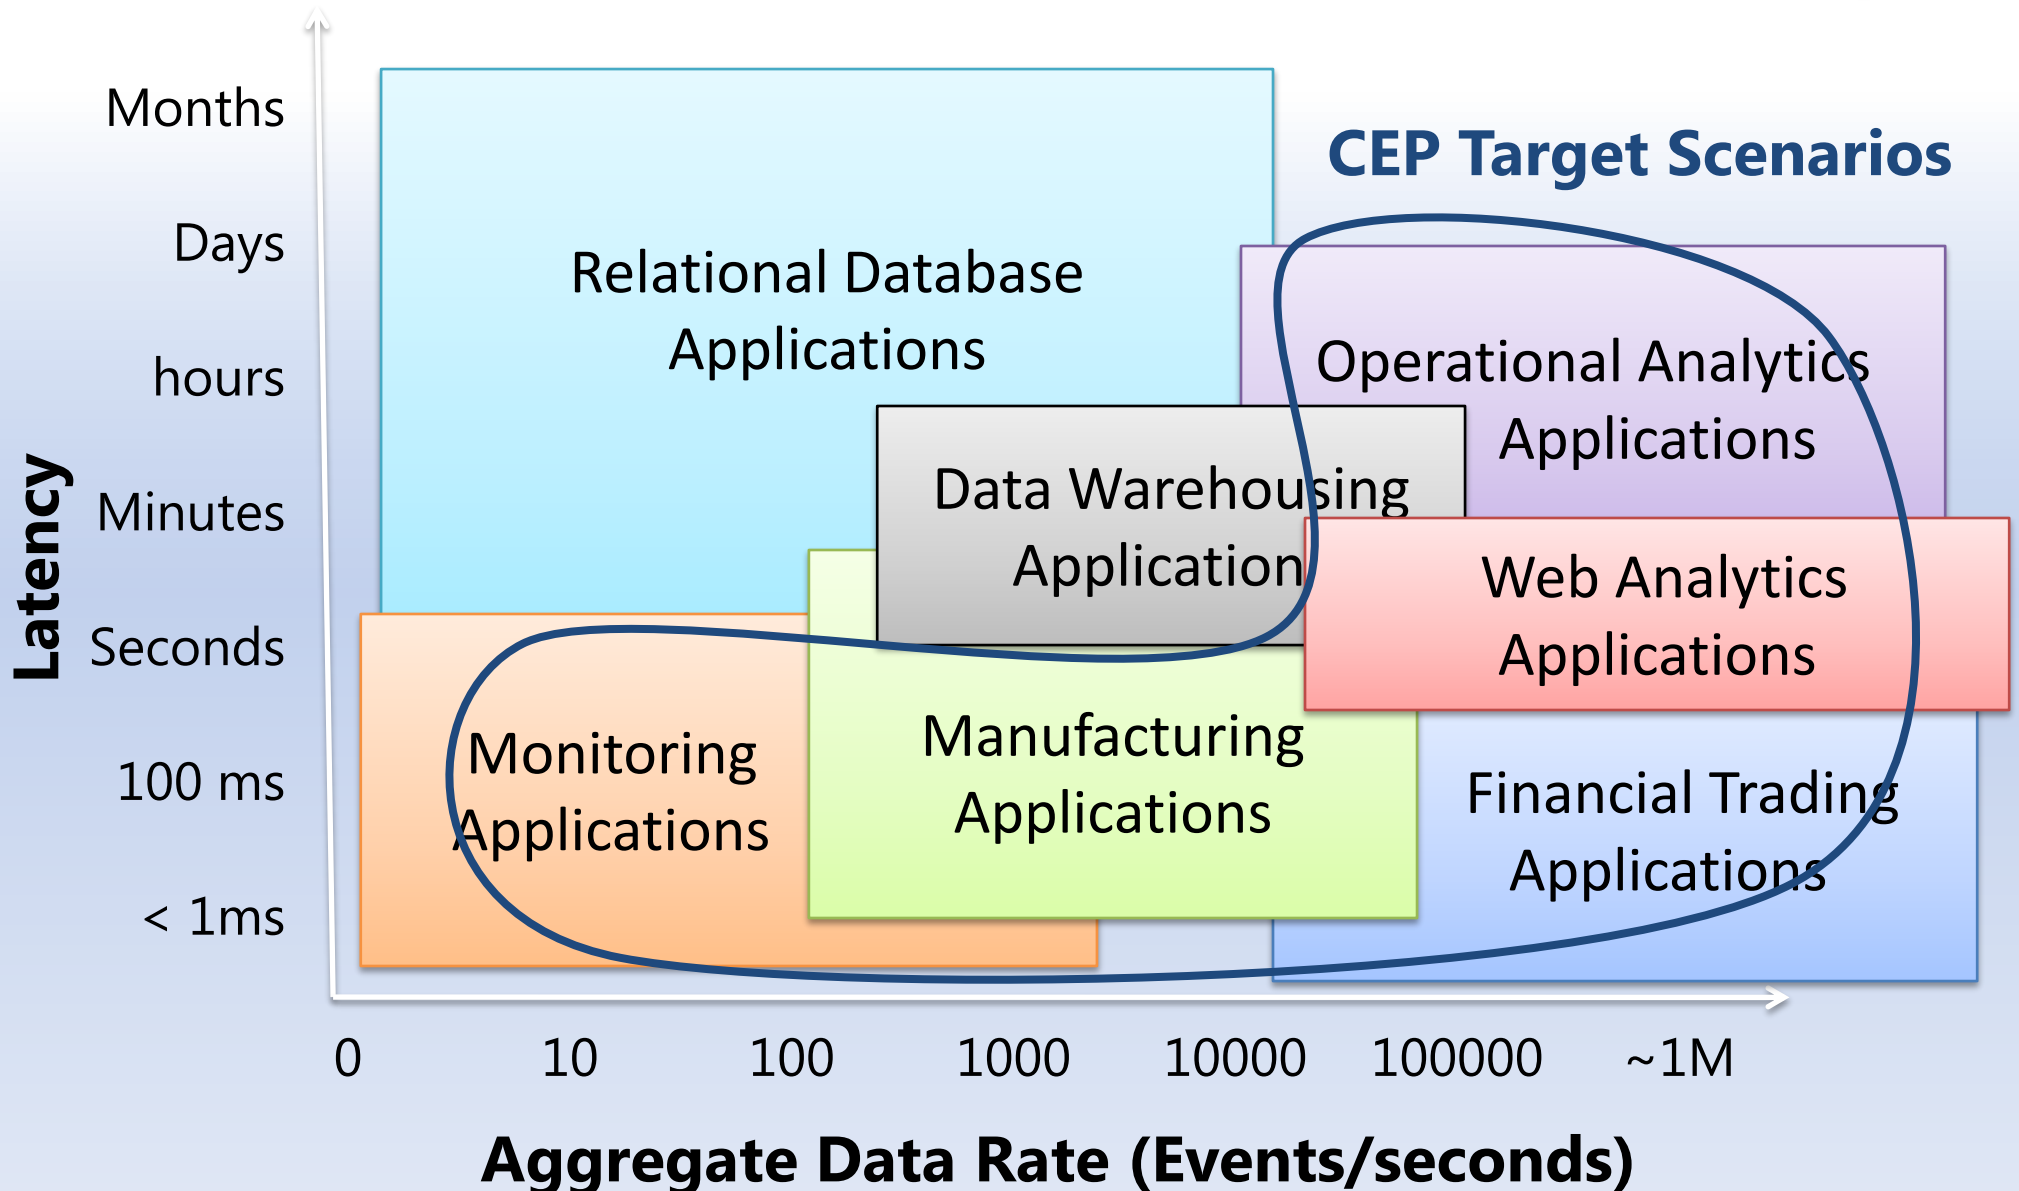
<http://www.flickr.com/photos/trainor/2902023575/>, Licensed under CC

## WHAT IS CEP ?

Complex Event Processing (CEP) is identifying meaningful patterns, relationships and data abstractions among unrelated events and fire an immediate response.

|                | Database Applications  | Event-driven Applications   |
|----------------|--|---|
| Query Paradigm | Ad-hoc queries or requests   | Continuous standing queries   |
| Latency        | Seconds, hours, days   | Milliseconds or less  |
| Data Rate      | Hundreds of events/sec   | Tens of thousands of events/sec or more   |
|                |  |  |

## SCENARIOS OF EVENT PROCESSING





## E-SCIENCE AND CEP



- E-Science often deals with National and Global scale usecases while we try to understand the world around us better.
- Following is an important class of E-science applications
  - Receives data about the world around us from sensors deployed across the country/world
  - Try to make sense, react to, predict, and/or control the world around us
  - Examples: Weather, Traffic Data, Surveillance, Smart Grid etc..
- CEP is a powerful enabling technologies for these usecases

## CEP QUERIES

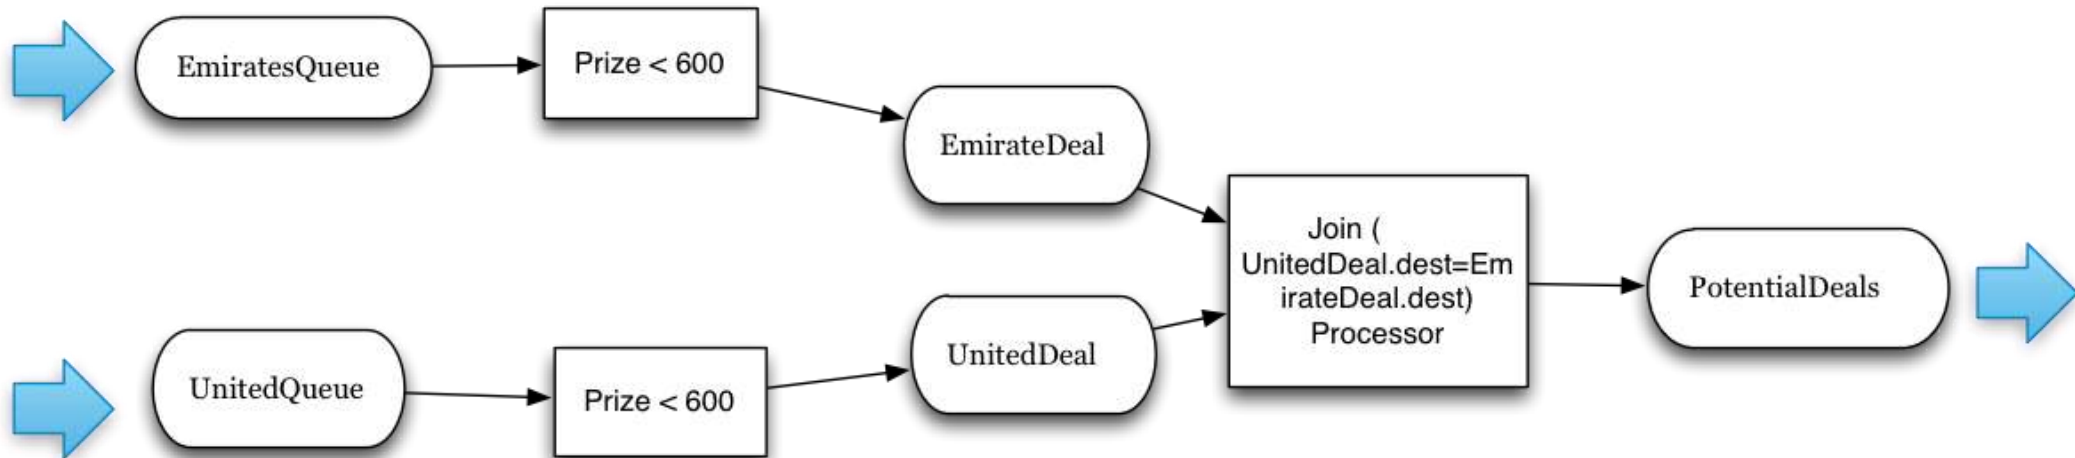
- Each Event consists of properties (name value pairs)
- We separate different events to streams
- Use a SQL like query language, but queries are evaluated on continuous event streams
- Types of queries are following
  - Selection (filtering) and projection (e.g. like select in SQL)
  - Windows – events are processed within a window (e.g. for aggregation and joins). Time window vs. length window.
  - Ordering – sequences and patterns (before, followed by conditions e.g. new location followed by small and a large purchase might suggest a fraud)
  - Join and split
  - Aggregation

## EXAMPLE QUERY

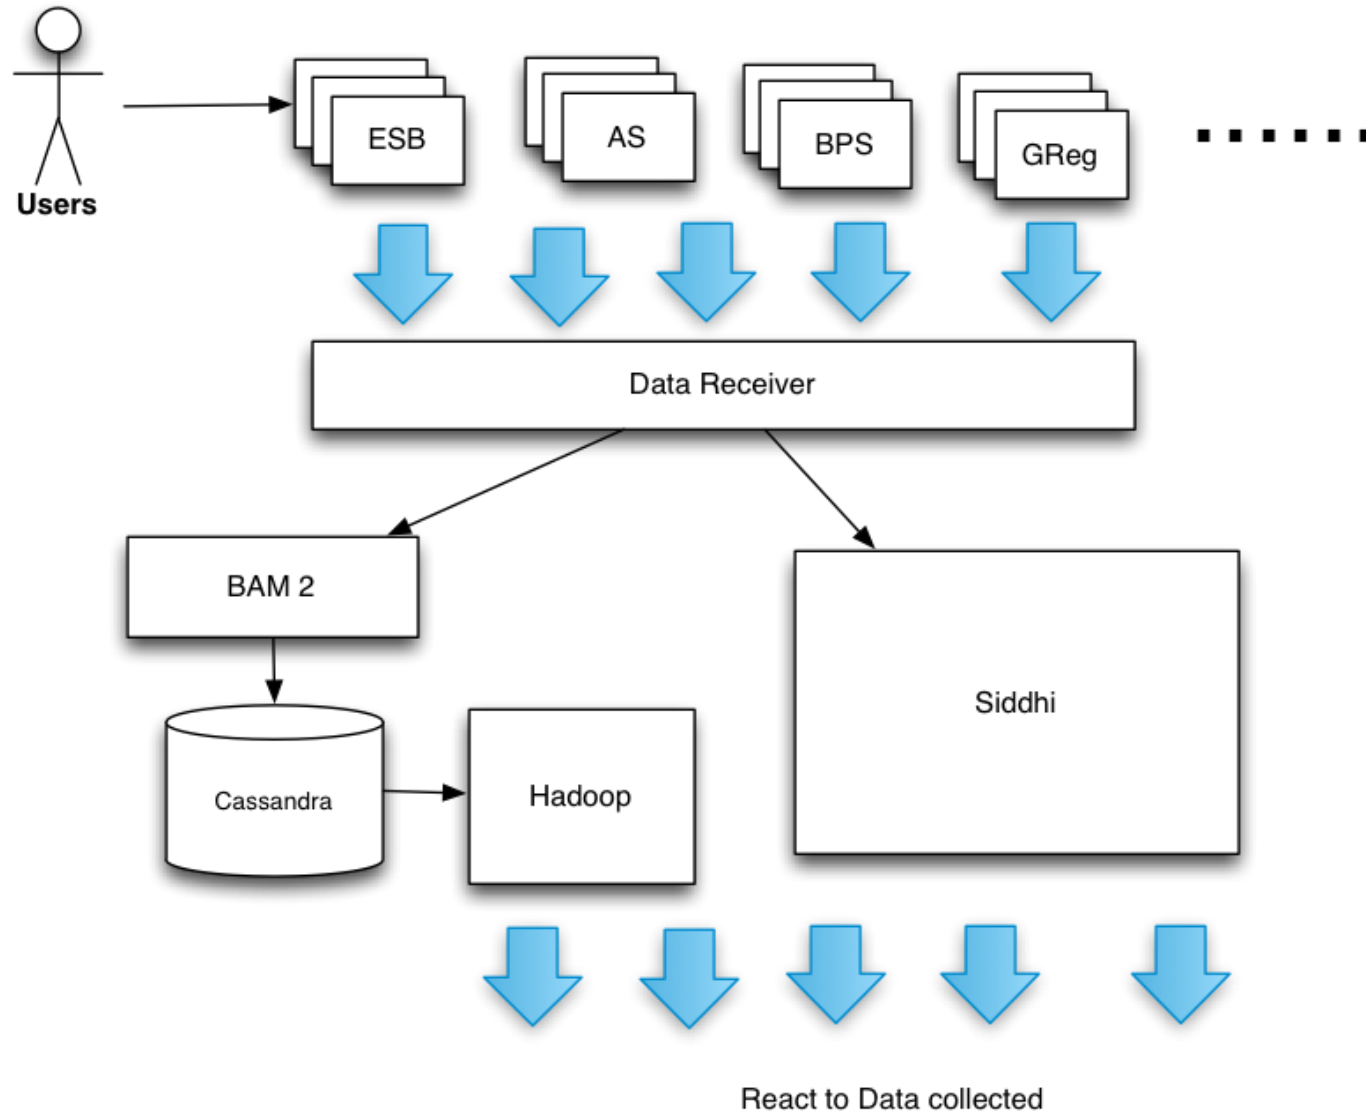
EmirateDeal= Select \* from Emirate where prize < 600

UnitedDeal = Select \* from United where prize < 600

PotentialDeal = Select \* from EmirateDeal join UnitedDeal  
where UnitedDeal .dest=EmirateDeal .dest

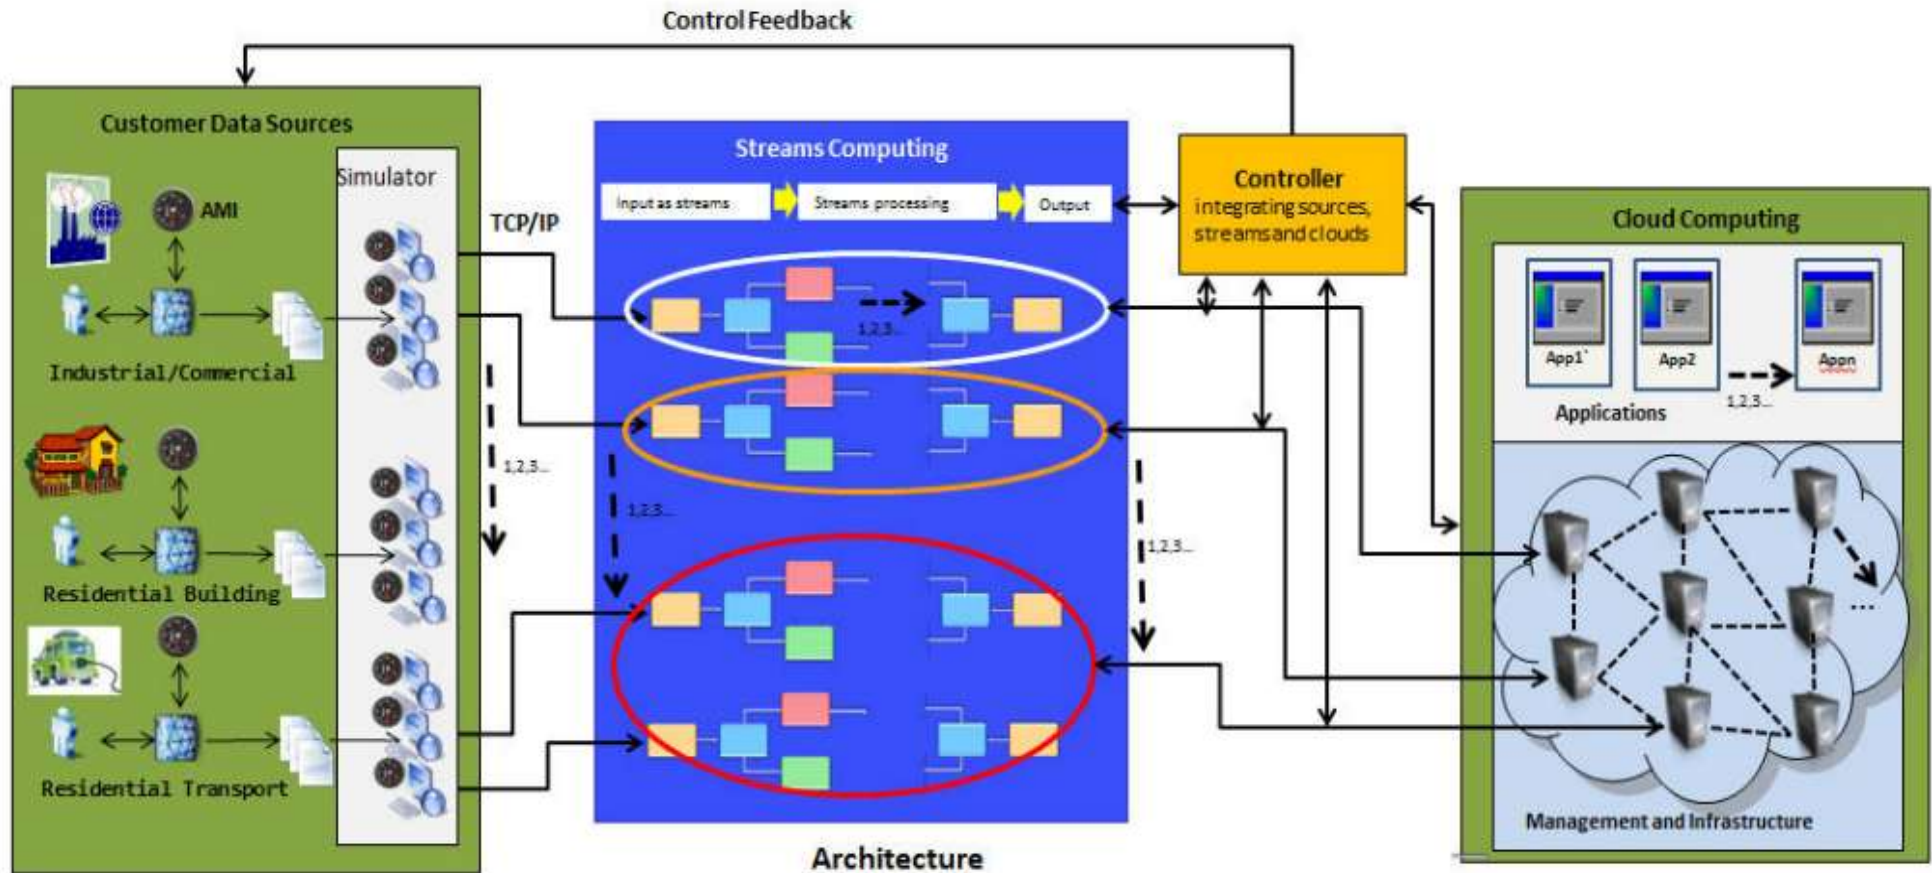


## CEP ROLE WITHIN WSO2 PLATFORM





## MOTIVATING USE CASE: LOS ANGELES SMART GRID DEMONSTRATION PROJECT



**Figure 2:** Design for adaptive rate control feedback from stream processing system to AMIs to optimize bandwidth into Cloud platform

## RELATED WORK ON CEP IMPLEMENTATIONS

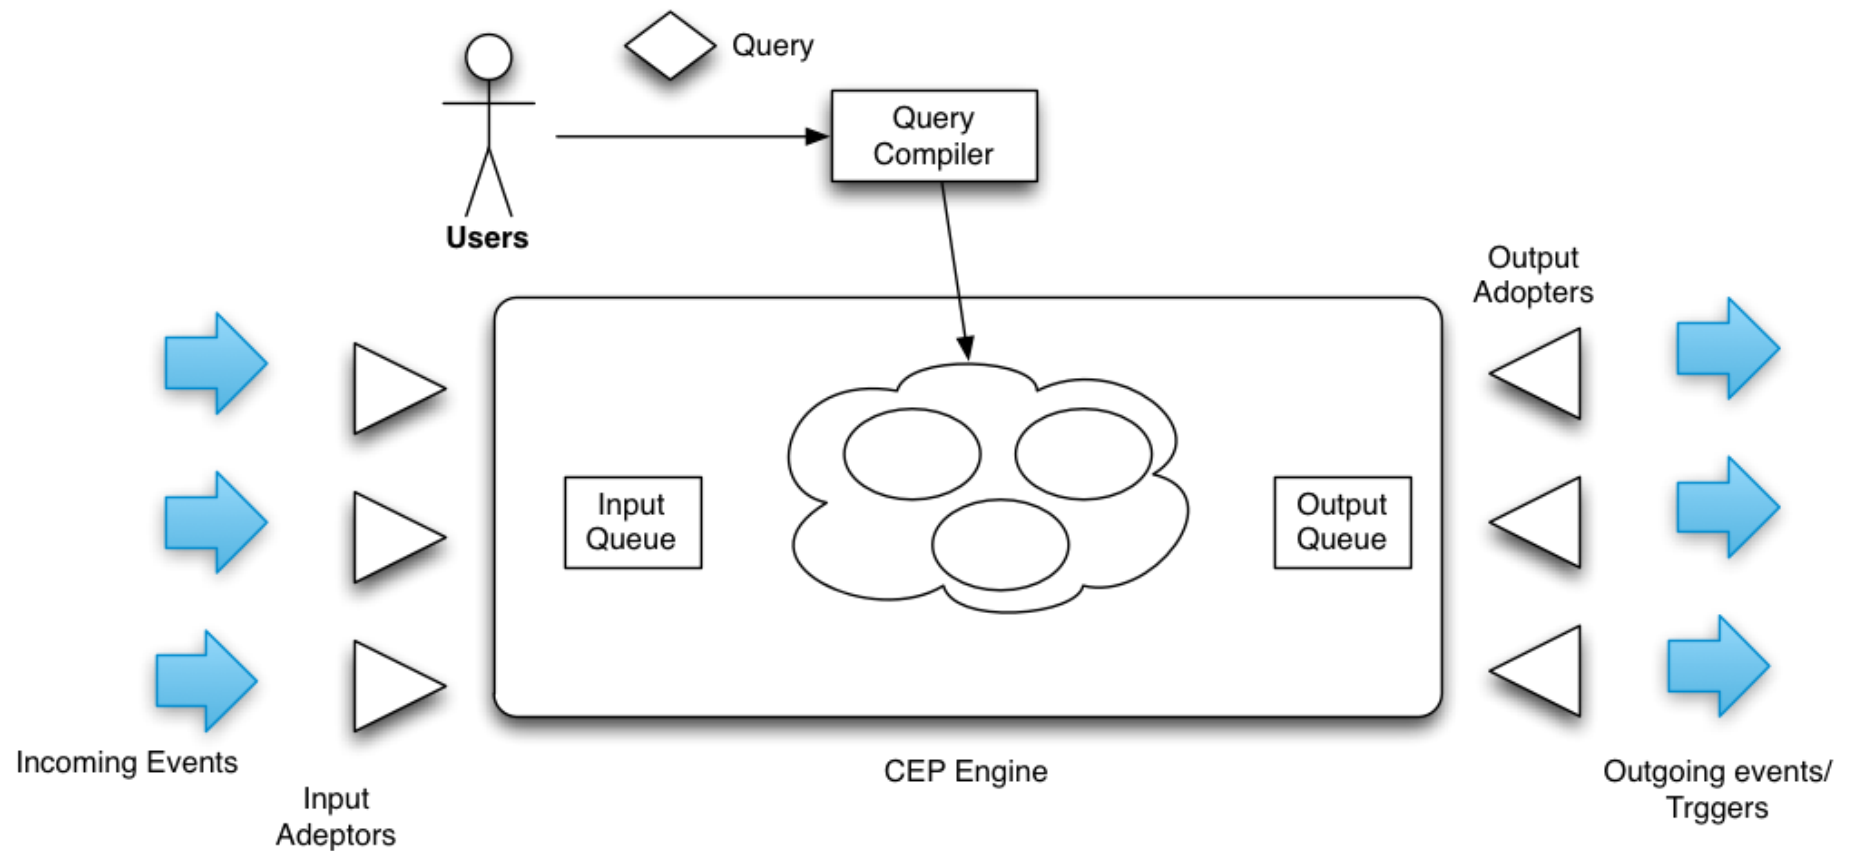
- A very good reference: “G. Cugola and A. Margara. Processing flows of information: From data stream to complex event” processing. ACM Computing Surveys, 2011.
- Initial work from Databases, active databases. But the “store and process” model was too slow.
- Stream processing uses a pub/sub model, very scalable but lacks temporal support
  - (Aurora [5, 2], PIPES [23, 7], STREAM [30], Borealis [6,1] S4 [25], and Storm.)
- CEP engines
  - Mainly uses a NFA based model
  - SASE (data flow and NFA)
  - Esper – NFA and Delta networks
  - Cayuga – NFA and re-subscriptions (single thread)

## GOAL: IDEAS FOR IMPROVEMENT

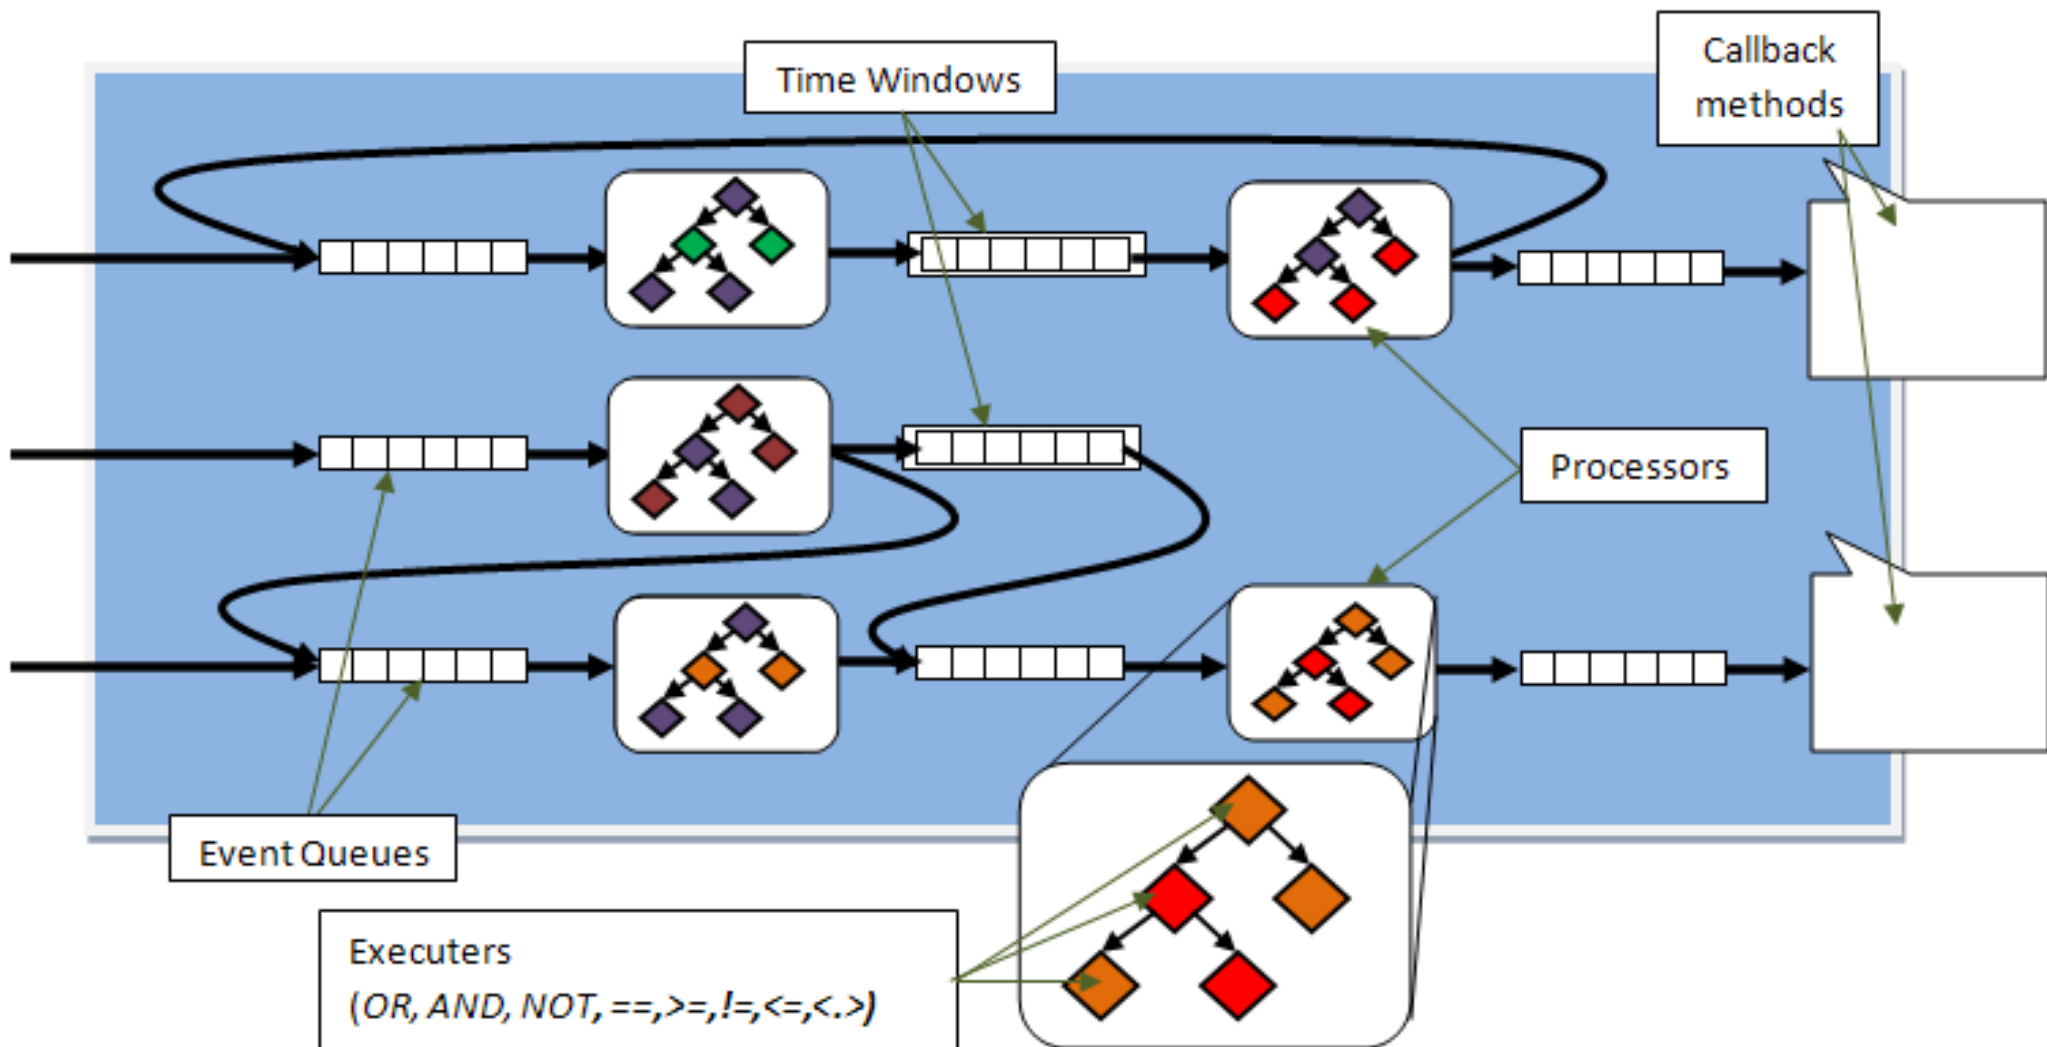


- Borrowing ideas from Stream processing: Pipeline based architecture with a pub/sub model inspired by Stream processing.
- Breaking single thread evaluations: Improving parallelism in processing through a pipeline based model
- A New State machine implementation that avoids checking a series of conditions
- Improved support for query chaining with complex queries

## SIDDHI HIGH LEVEL ARCHITECTURE



## SIDDHI PROCESSING IMPLEMENTATION





## SIDDHI PROCESSING IMPLEMENTATION

- Uses a pipeline architecture
- Consist of a pipeline of processors (support all other operators except windows)
- Processors are connected by queues (they implement windows)
- Processors take inputs from a queue by subscribing to them and placing results in a output queue
- Many threads are assigned to processors, where they take data from input queues, process them, and place them in output queues.

## API TO WRITE QUERIES

```
//Instantiate SiddhiManager
SiddhiManager siddhiManager = new SiddhiManager();
//Get the QueryFactory to create queries
QueryFactory qf = siddhiManager.getQueryFactory();

InputEventStream stockStream = new InputEventStream(
    "StockStream",
    new String[]{"symbol", "price", "volume"},
    new Class[]{String.class, Float.class, Long.class}
);
//Setting a time window of 1 hour
stockStream.setTimeWindow(3600000);
//Assign input stream
siddhiManager.addInputEventStream(stockStream);

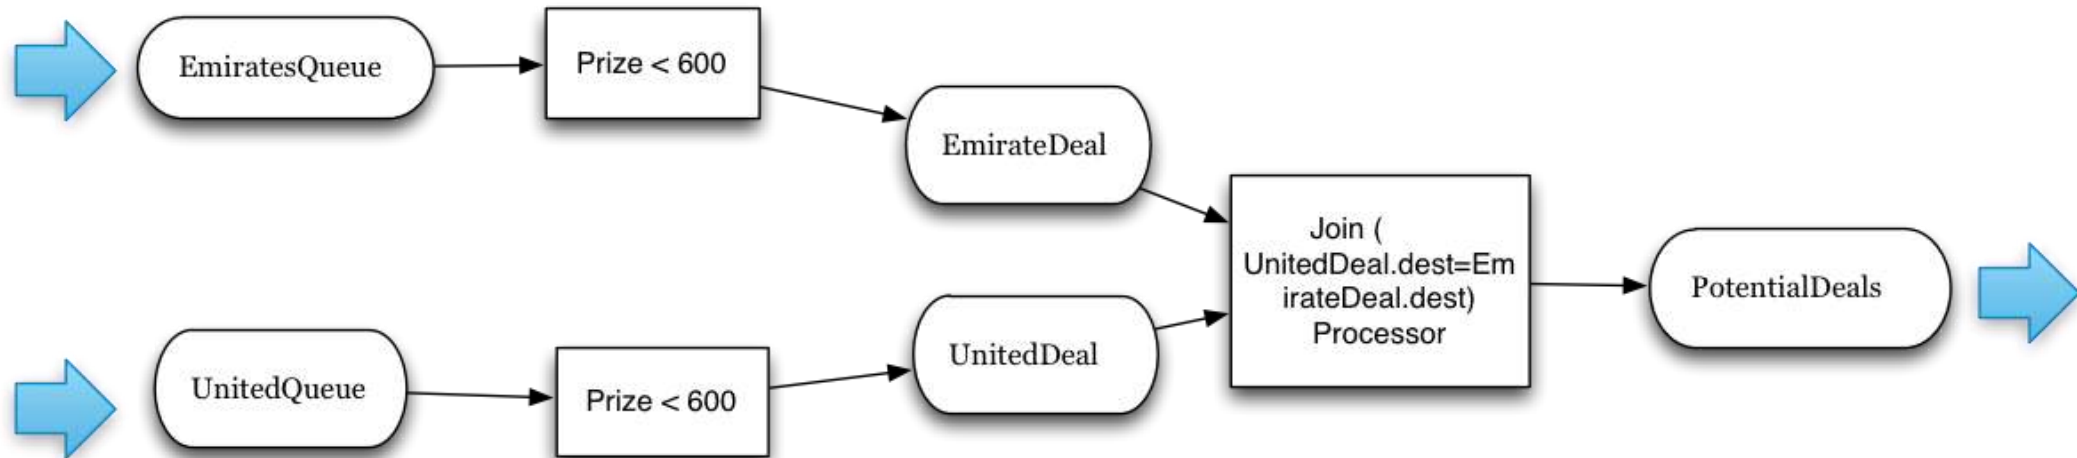
//Create Query
Query query = qf.createQuery(
    "StockQuote",
    qf.output("symbol=StockStream.symbol", "price=avg(StockStream.price)"),
    qf.inputStream(stockStream),
    qf.and(
        qf.condition("StockStream.symbol", EQUAL, "IBM"),
        qf.condition("StockStream.volume", GREATERTHAN, "10000")
    )
);
//Assign query
siddhiManager.addQuery(query);
```

## EXAMPLE QUERY IMPLEMENTATION

EmirateDeal= Select \* from Emirate where prize < 600

UnitedDeal = Select \* from United where prize < 600

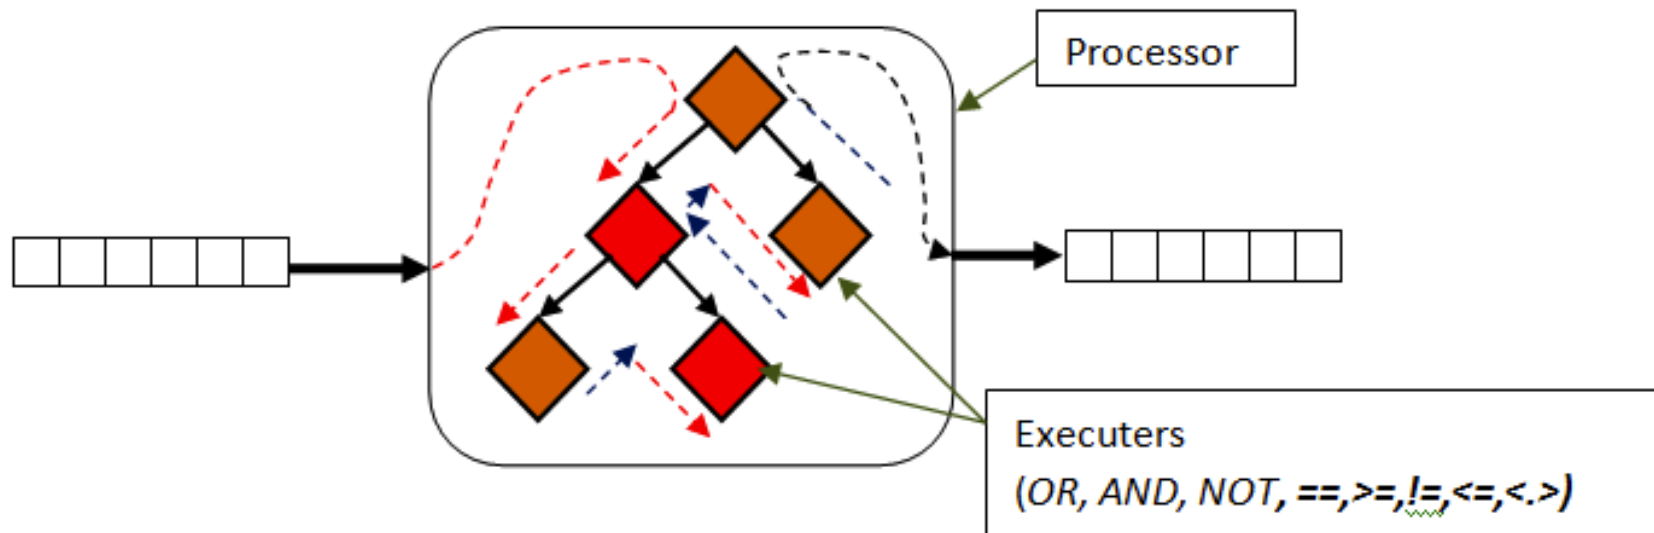
PotentialDeal = Select \* from EmirateDeal join UnitedDeal  
where UnitedDeal .dest=EmirateDeal .dest



## IMPLEMENTING FILTERS


```
Query query = qf.createQuery(  
    "StockQuote",  
    qf.output("symbol=StockStream.symbol","price=avg(StockStream.price)"),  
    qf.inputStream(stockStream),  
    qf.and(  
        qf.condition("StockStream.symbol", EQUAL, "IBH"),  
        qf.condition("StockStream.volume", GREATERTHAN, "10000")  
    )  
);
```

- Evaluate the tree, and optimizations to stop the evaluations as soon as possible



## IMPLEMENTING WINDOWS

- Implemented as a queue
- Use events to notify the windows event addition and removals using a pub/sub pattern to processors

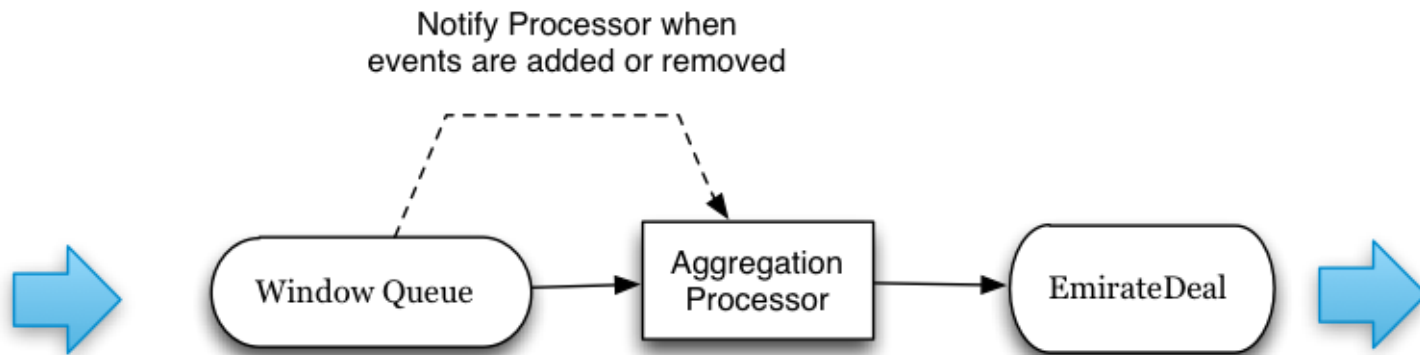




## AGGREGATION


- Siddhi supports avg, sum, count, max, min

```
Query query = qf.createQuery(  
    "StockQuote",  
    qf.output("symbol=CSEStream.symbol", "avgPrice=avg(CSEStream.price)"),  
    qf.inputStream(cseEventStream),  
    qf.condition("CSEStream.price", GREATERTHAN_EQUAL, "20")  
);  
query.groupBy("CSEStream.symbol");  
query.having(qf.condition("StockQuote.avgPrice", GREATERTHAN, "50"));
```



- These are implemented within processor and each addition or removal of a event updates the value (if it is a window based, queues notify subscribers when thing change).

## IMPLEMENTING JOINS



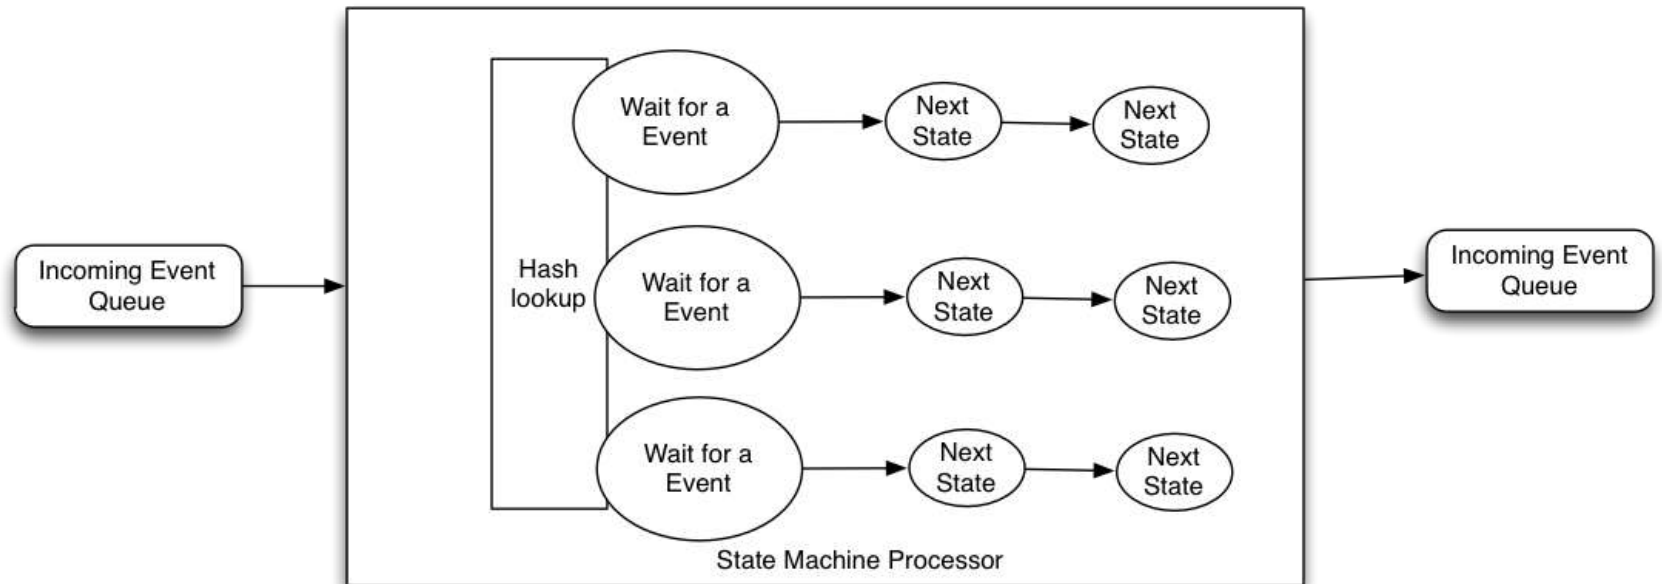
- Implemented as a processor that moves data from several queues to a different queue after joining
- When an event arrives, if it matches the join condition, we send them over to the next queue, else keeps them to match with the future events.
- We keep them as long as the window condition allows (this is done by keeping references to queues, and checking within them for matching )

## PATTERNS AND SEQUENCES

- Patterns and sequences handle series of events like A->B->C, there are two classes patterns and sequences
- Patterns matches with events ignoring events in between while sequences matches exactly. You can use star (\*) with sequences to mimic patterns
- Supported through a state machine (NFA)

```
Query query = qf.createQuery(  
    "StockQuote",  
    qf.output("priceA=$0.first.price", "priceA=$0.last.price", "priceB=$1.price"),  
    qf.inputStreams(cseEventStream, infoStock),  
    qf.sequence(  
        qf.star(  
            qf.condition("CSEStream.price", GREATERTHAN, "$0.prev.price")  
        ),  
        qf.condition("CSEStream.price", LESSTHAN, "$0.last.price")  
    )  
);
```

## IMPLEMENTING PATTERNS AND SEQUENCES



- Use a event model
- Say we need A->B->C, then we break down it to list of listeners, which when triggered remove itself and registers the next in the list.
- We apply different variations of this to support \*, and sequences .

## PERFORMANCE RESULTS




- We compared Siddhi with Esper, the widely used opensource CEP engine
- For evaluation, we did setup different queries using both systems, push events in to the system, and measure the time till all of them are processed.
- We used Intel(R) Xeon(R) X3440 @2.53GHz , 4 cores 8M cache 8GB RAM running Debian 2.6.32-5-amd64 Kernel



## PERFORMANCE COMPARISON WITH ESPER

Simple filter without window


```
select symbol, price  
from StockTick(price>6)
```



## PERFORMANCE COMPARISON WITH ESPER

Average calculation of a given symbol within time window


```
select iirstream symbol, price, avg(price)
from StockTick(symbol='IBM').win:time(.005)
```



## PERFORMANCE COMPARISON WITH ESPER

State machine query for pattern matching

```
select f.symbol, p.accountNumber, f.accountNumber  
from pattern [every f=FraudWarningEvent2 -> p=PINChangeEvent2(accountNumber  
= f.accountNumber)]
```



## **SIDDHI vs. ESPER (FUNCTIONALITY)**

- Following describe Siddhi in contrast to Esper using the comparisons defined in Cugola et al.
- In terms of Functional and Processing Models and Data, Time, and Rule Models (see Table I and III of Cugola et al.), Esper and Siddhi behave the same.
- In terms of Supported language model (table IV),
  - Esper supports Pane and Tumble windows, User Defined operators and Parameterization while Siddhi does not support them
  - Siddhi supports removal of duplicates that is not supported by Siddhi.
- Both support all other language constructs.

## FEATURES WE DID NOT DISCUSSED

- Improvements in the Selection Operator
  - using MVEL
- Support conditions on the aggregation functions
  - having, group-by
- Event quantifications in sequential patterns
  - kleene closure (\*)
- Unique function



## CURRENT USERS OF SIDDHI

- "Los Angeles Smart Grid Demonstration Project"
  - It forecasts electricity demand, respond to peak load events, and improves sustainable use of energy by consumers.  
<http://ceng.usc.edu/~simmhan/pubs/simmhan-usctr2011-smartgridinformatics.pdf>
- Open MRS NCD module — idea is to detect and notify patient when certain conditions have occurred <https://wiki.openmrs.org/display/docs/Notifiable+Condition+Detector+%28NCD%29+Module>
- WSO2 CEP Server (Perspective)

## FUTURE WORK



- SQL like query language for Siddhi
- Building scalable processing network using Siddhi nodes
  - Siddhi is currently a java library, build a Thrift based server for Siddhi
  - Supporting a network of Siddhi processors (clustering)
  - Query partition and optimization for scale
- Support for Pane and Tumble windows, User Defined operators and Parameterization
- Available under Apache License

## IMMEDIATE PLANS



- Have resource allocations and development efforts to continue next year
- Looking forward to build a opensource community around it
- Looking for usecases, real data feeds, and of course Guinea pigs

# *Siddhi-CETP*

**QUESTIONS?**

